- Part 2

[slide 10]

(you can copy the files from part 1 to another folder, and start from there)

Let's build a web app.

Go to http://flask.pocoo.org/!

```
$ pip install Flask
```

Change `app.py` to:

```python
from flask import Flask
app = Flask(__name__)

@app.route("/")
def hello():
    return "Hello World!"

app.run(host="0.0.0.0", debug=True)
```

Run the app and check http://0.0.0.0:5000/:

```
$ python app.py
```

Remove the `sleeper` from `docker-compose.yml`:

```yaml
version: "3"
services:
  app:
    build: .
```

Run `docker-compose up --build` and check http://0.0.0.0:5000/.

It doesn't work. Why?

We need to expose and publish the container's port `5000` to the host (our machine):

```yaml
version: "3"
services:
  app:
    build: .
    ports:
      - "3333:5000"
```

The above publishes container's port `5000` on host's port `3333`.

Run `docker-compose up` and http://0.0.0.0:3333/.

Deploy it on Kubernetes

[slide 11]

Create file `app.yml` (simply based on Kubernetes 101):

```yaml
apiVersion: v1
kind: Pod
metadata:
  name: app
  labels:
    foo: vitor
spec:
  containers:
  - name: app
    image: vitorenesduarte/tutorial
```

(Compared to 101, we added `foo: vitor` as a label, because `kubectl expose`, which we will use later, requires it)

And deploy it on Kubernetes.

(for that you need `CONFIG`, a Kubernetes configuration file, which I will provide)

```
$ kubectl --kubeconfig=CONFIG create -f app.yml
```

QUESTION: will this work?

```
$ kubectl --kubeconfig=CONFIG get pods
NAME        READY       STATUS          RESTARTS    AGE
app         0/1         ErrImagePull    0           4s
```

Ups. The docker image is still local.

Let's push it to Docker Hub.

Create an account there, and login with `docker login`. Then:

```
$ docker build -t vitorenesduarte/tutorial .
$ docker push vitorenesduarte/tutorial
```

Before anything else, let's avoid always having to specify `--kubeconfig`.

Let's check the manual.

```
$ kubectl config --help | sed -n '5,7p'
```

We can simply have `$KUBECONFIG` environment variable pointing to the `CONFIG` file, e.g.:

```
$ export KUBECONFIG=$(pwd)/CONFIG
```

Now, let's delete the `app` pod and deploy again.

```
$ kubectl delete pod app
$ kubectl get pods
$ kubectl create -f app.yml
$ kubectl get pods --watch
```

Create a load balancer so that we can access our app:

Create file `app-service.yml`:

```yaml
apiVersion: v1
kind: Service
metadata:
  name: app-service
spec:
  type: LoadBalancer
  ports:
  - port: 5000
    targetPort: 3333
  selector:
    foo: vitor
```

```
$ kubectl create -f app-service.yml
$ kubectl get service app-service
TODO
```

Watch until `EXTERNAL-IP` is no longer ':

```
$ kubectl get service app-service --watch
```

And then go to http://EXTERNAL-IP:3333:

Does the load balancing work?

Let's slightly change our app, so that each pod has an identifier.

```python
from flask import Flask
app = Flask(__name__)

import sys
id = sys.argv[1] if len(sys.argv) > 1 else "ups!"
```

```python
@app.route("/")
def hello():
    return "Hello World! (from " + id + ")"

app.run(host="0.0.0.0", debug=True)
```

Change the `Dockerfile`, so that we can pass the pod identifier as an environment variable `$ID`:

```dockerfile
FROM python:alpine

RUN pip install flask

COPY app.py /

CMD python app.py $ID
```

As identifier, we'll use the pod name.

We can change `app.yml` so that we pass the the pod name as environment variable `$ID`:

```yaml
apiVersion: v1
kind: Pod
metadata:
  name: app-1
  labels:
    foo: vitor
spec:
  containers:
  - name: app
    image: vitorenesduarte/tutorial
    env:
    - name: ID
      valueFrom:
        fieldRef:
          fieldPath: metadata.name
---
apiVersion: v1
kind: Pod
metadata:
  name: app-2
  labels:
    foo: vitor
spec:
  containers:
  - name: app
    image: vitorenesduarte/tutorial
    env:
    - name: ID
      valueFrom:
        fieldRef:
          fieldPath: metadata.name
```

Now go to http://EXTERNAL-IP:3333, and see the identifier changing.